

# Processing the Data with MapReduce

Prof. Abhaya Kumar Sahoo  
School of Computer Engineering  
Kalinga Institute of Industrial Technology  
Deemed to be University, Bhubaneswar

# Contents

- The introduction of MapReduce,
- MapReduce Architecture,
- Data flow in MapReduce Splits,
- Mapper,
- Portioning,
- Sort and shuffle,
- Combiner,
- Reducer,
- Basic Configuration of MapReduce,
- MapReduce life cycle,
- Driver Code,
- Mapper and Reducer,
- How MapReduce Works.

# The introduction of MapReduce

- MapReduce is a processing technique and a program model for distributed computing based on java.
- contains two important tasks, namely Map and Reduce.
- **Map** takes a set of data and converts it into another set of data, where individual elements are broken down into tuples (key/value pairs).
- The processing primitive is called mapper.
- **Reduce** task takes the output from a map as an input and combines those data tuples into a smaller set of tuples.
- Reduce task is always performed after the map job.
- The processing primitive is called reducer.
- **The major advantage of MapReduce is that it is easy to scale data processing over multiple computing nodes.**

# The introduction of MapReduce

➤ The main advantage is that we write an application in the MapReduce form, scaling the application to run over hundreds, thousands, or even tens of thousands of machines in a cluster with a configuration change.

- MapReduce program executes in three stages:

- ❖ map stage,
- ❖ shuffle stage,
- ❖ reduce stage.

➤ **Map stage** : The map or mapper's job is to process the input data.

➤ Generally the input data is in the form of file or directory and is stored in the Hadoop file system (HDFS).


➤ The input file is passed to the mapper function line by line.

➤ The mapper processes the data and creates several small chunks of data.

## The introduction of MapReduce

- **Reduce stage** : This stage is the combination of the **Shuffle** stage and the **Reduce** stage.
- The Reducer's job is to process the data that comes from the mapper.
- After processing, it produces a new set of output, which will be stored in the HDFS.

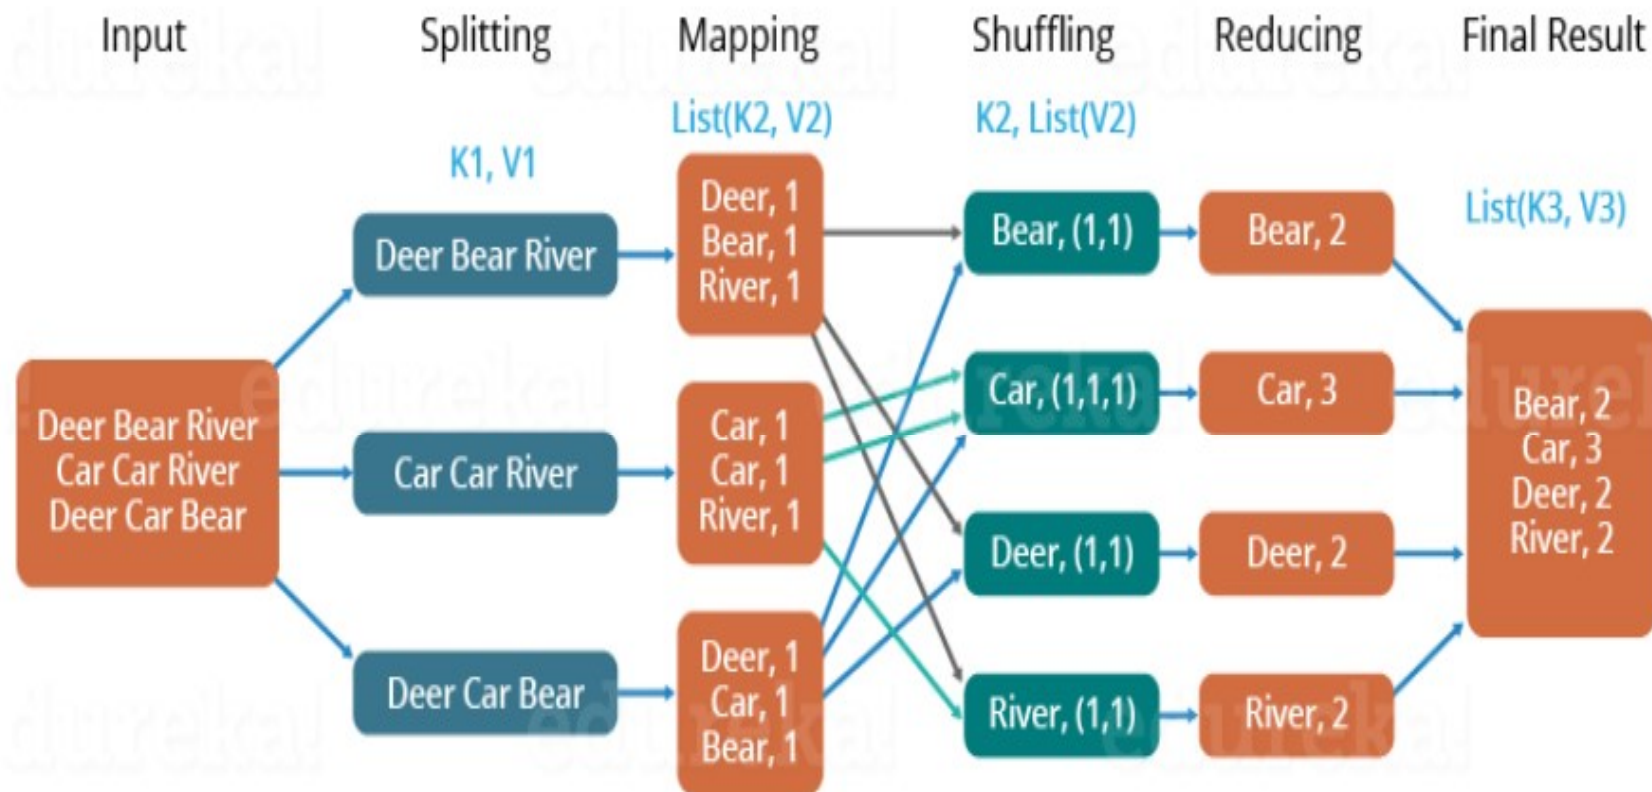
## Fundamental Principle



- ❖ During a MapReduce job, Hadoop sends the Map and Reduce tasks to the appropriate servers in the cluster.
- ❖ The framework manages all the details of data-passing such as issuing tasks, verifying task completion, and copying data around the cluster between the nodes.


## Fundamental Principle

- ❖ Most of the computing takes place on nodes with data on local disks that reduces the network traffic.
- ❖ After completion of the given tasks, the cluster collects and reduces the data to form an appropriate result, and sends it back to the Hadoop server.





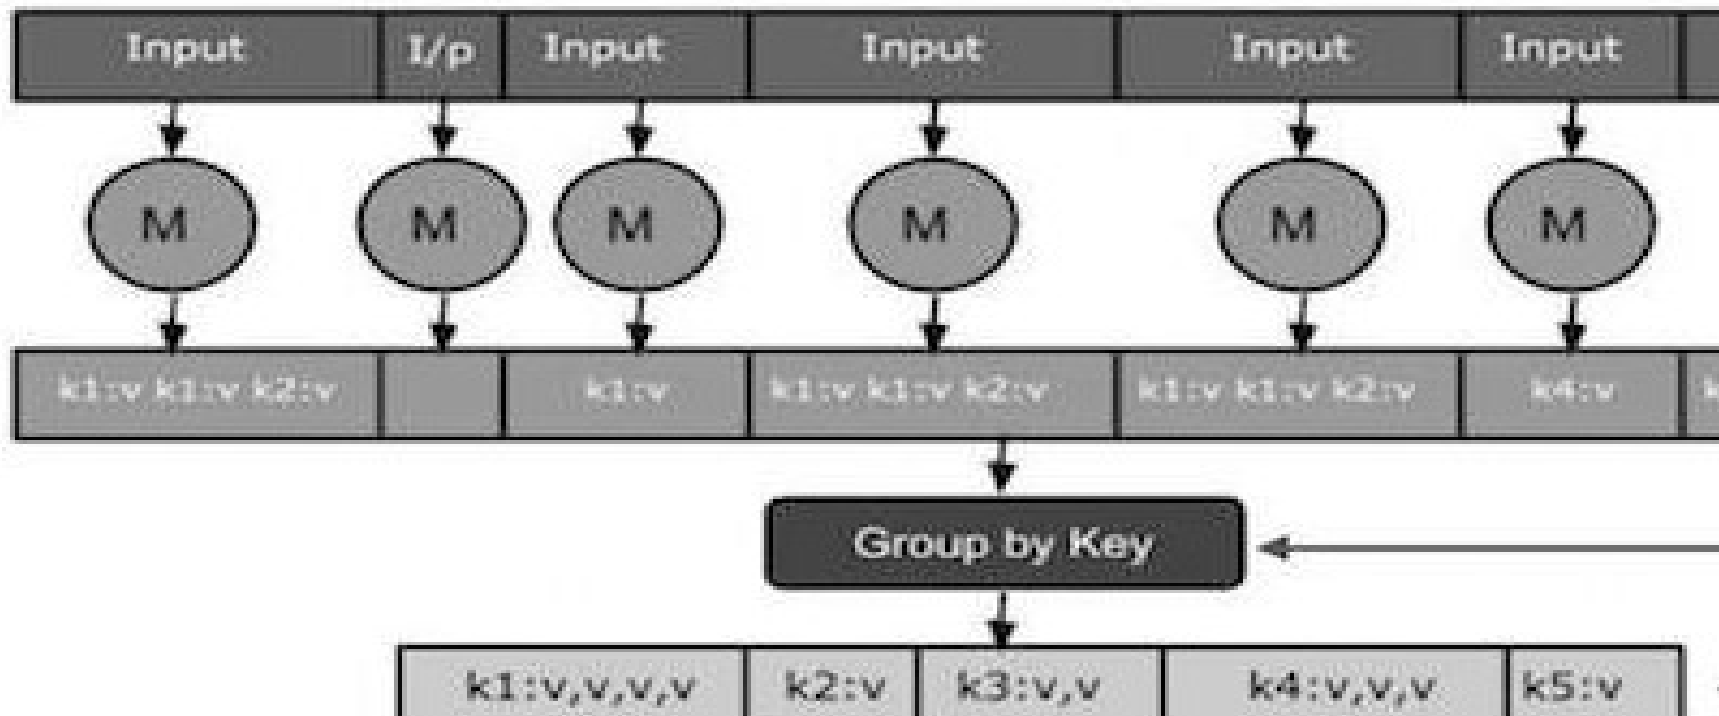
## Why MapReduce?



❖ MapReduce divides a task into small parts and assigns them to many computers. Later, the results are collected at one place and integrated to form the result dataset.

## MapReduce Works Principle

- ❖ The Map task takes a set of data and converts it into another set of data, where individual elements are broken down into tuples (key-value pairs).
- ❖ The Reduce task takes the output from the Map as an input and combines those data tuples (key-value pairs) into a smaller set of tuples.



## Terms

**Input Phase** – Here we have a Record Reader that translates each record in an input file and sends the parsed data to the mapper in the form of key-value pairs.

**Map** – Map is a user-defined function, which takes a series of key-value pairs and processes each one of them to generate zero or more key-value pairs.

**Intermediate Keys** – The key-value pairs generated by the mapper are known as intermediate keys.

**Combiner** – A combiner is a type of local Reducer that groups similar data from the map phase into identifiable sets.

- It takes the intermediate keys from the mapper as input and applies a user-defined code to aggregate the values in a small scope of one mapper.
- It is not a part of the main MapReduce algorithm, it is optional.

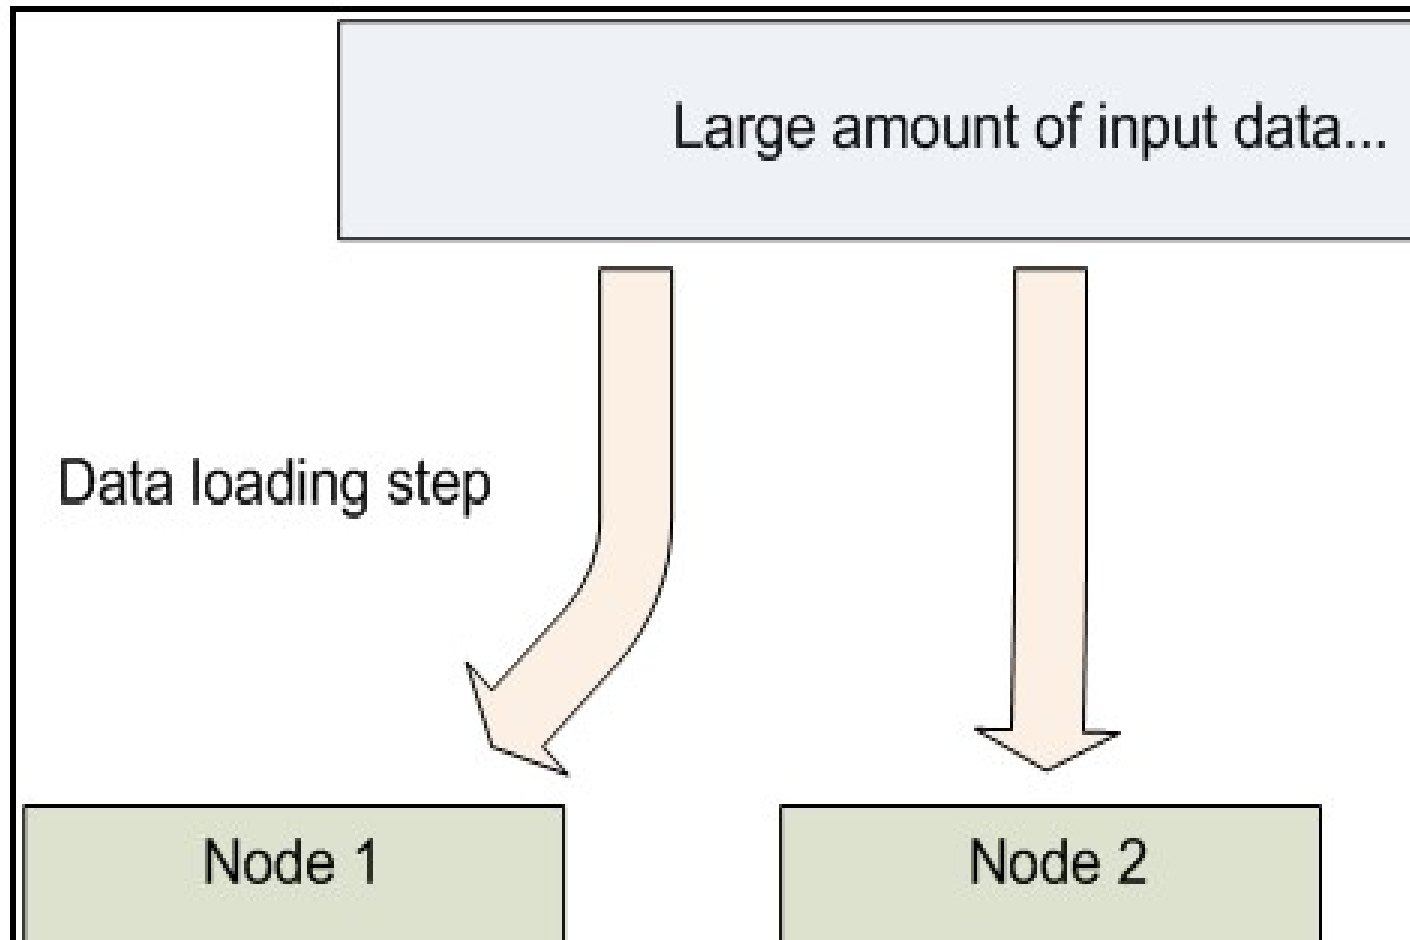
## Terms

**Shuffle and Sort** – The Reducer task starts with the Shuffle and Sort step. It downloads the grouped key-value pairs onto the local machine, where the Reducer is running. The individual key-value pairs are sorted by key into a larger data list. The data list groups the equivalent keys together so that their values can be iterated easily in the Reducer task.


**Reducer** – The Reducer takes the grouped key-value paired data as input and runs a Reducer function on each one of them. Here, the data can be aggregated, filtered, and combined in a number of ways, and it requires a wide range of processing. Once the execution is over, it gives zero or more key-value pairs to the final step.

**Output Phase** – In the output phase, we have an output formatter that translates the final key-value pairs from the Reducer function and writes them onto a file using a record writer.

# Data Distribution

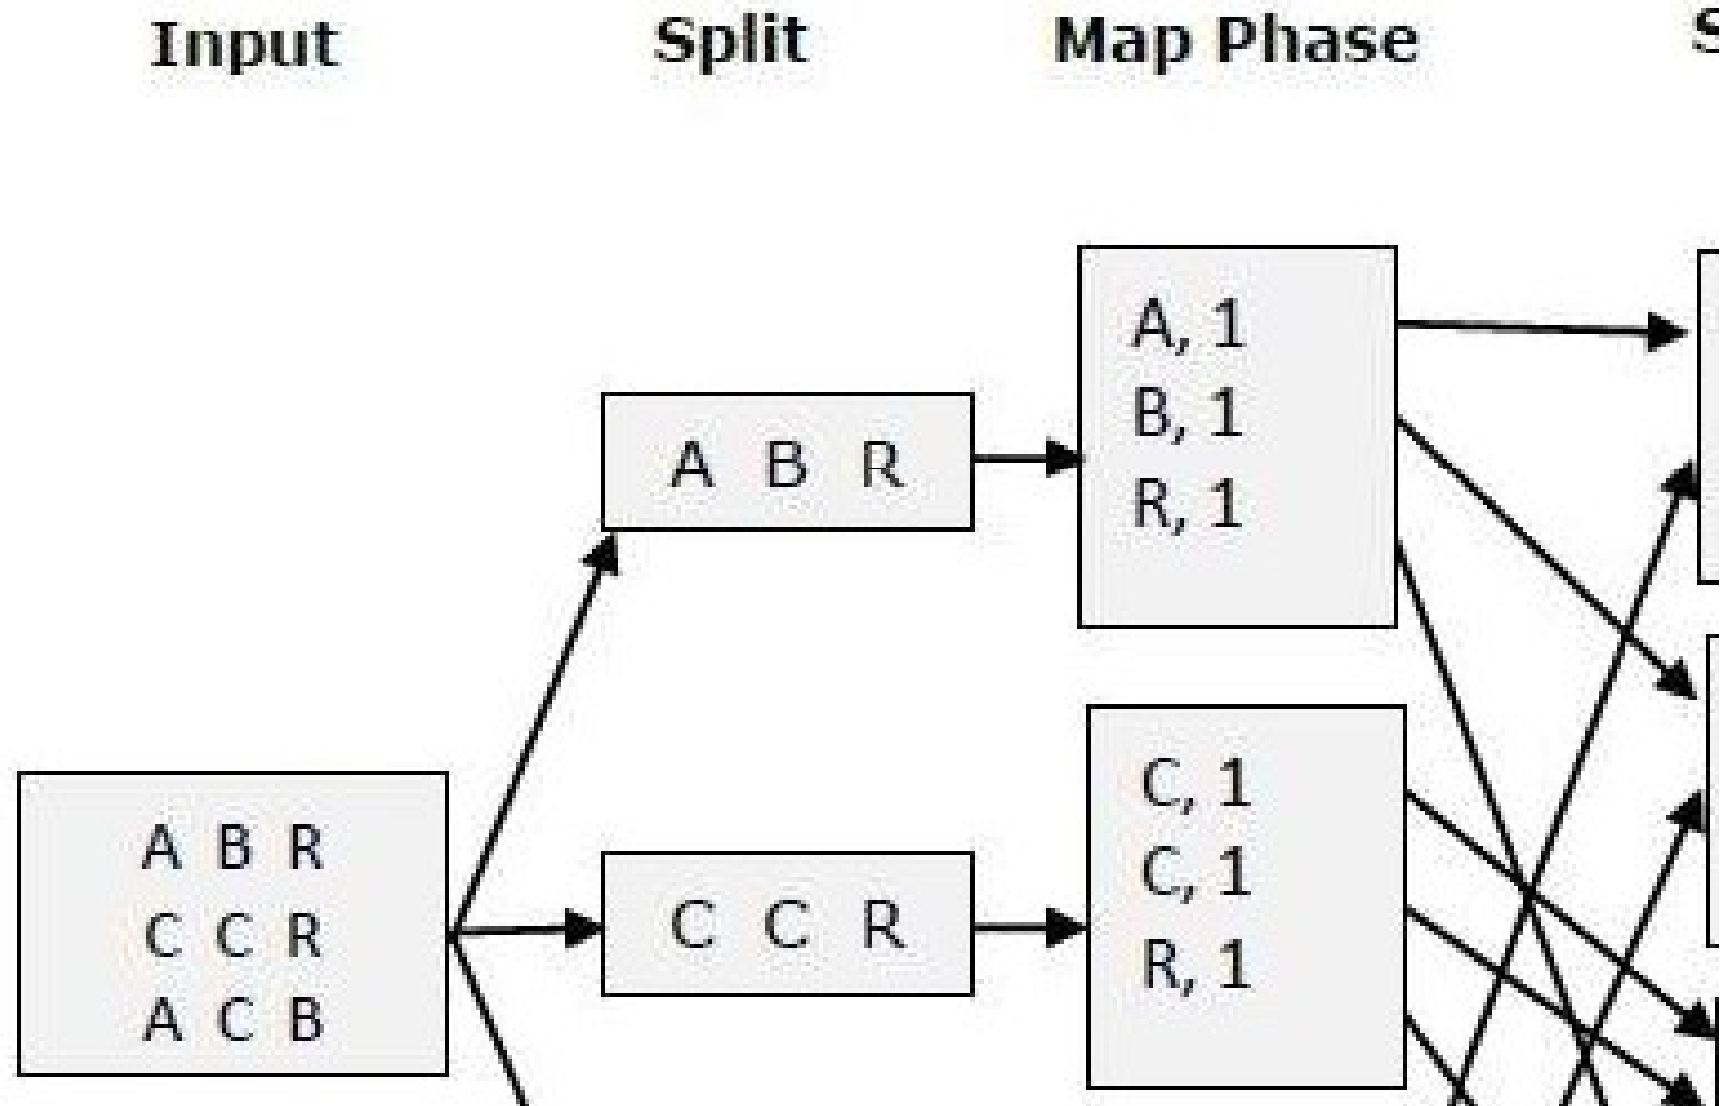


# MapReduce: Isolated Processes



In MapReduce, records are processed in isolation by tasks called *Mappers*. The output from the Mappers is then brought together into a second set of tasks called *Reducers*, where results from different mappers can be merged together.

## Simple Example



## Terms

**Shuffle and Sort** – The Reducer task starts with the Shuffle and Sort step. It downloads the grouped key-value pairs onto the local machine, where the Reducer is running. The individual key-value pairs are sorted by key into a larger data list. The data list groups the equivalent keys together so that their values can be iterated easily in the Reducer task.

**Reducer** – The Reducer takes the grouped key-value paired data as input and runs a Reducer function on each one of them. Here, the data can be aggregated, filtered, and combined in a number of ways, and it requires a wide range of processing. Once the execution is over, it gives zero or more key-value pairs to the final step.

**Output Phase** – In the output phase, we have an output formatter that translates the final key-value pairs from the Reducer function and writes them onto a file using a record writer.



## Programming Terminology

**PayLoad** - Applications implement the Map and the Reduce functions, and form the core of the job.

**Mapper** - Mapper maps the input key/value pairs to a set of intermediate key/value pair.

**NamedNode** - Node that manages the Hadoop Distributed File System (HDFS).

**DataNode** - Node where data is presented in advance before any processing takes place.

**MasterNode** - Node where JobTracker runs and which accepts job requests from clients.

**SlaveNode** - Node where Map and Reduce program runs.

**JobTracker** - Schedules jobs and tracks the assign jobs to Task tracker.

## Programming Terminology

**Task Tracker** - Tracks the task and reports status to JobTracker.

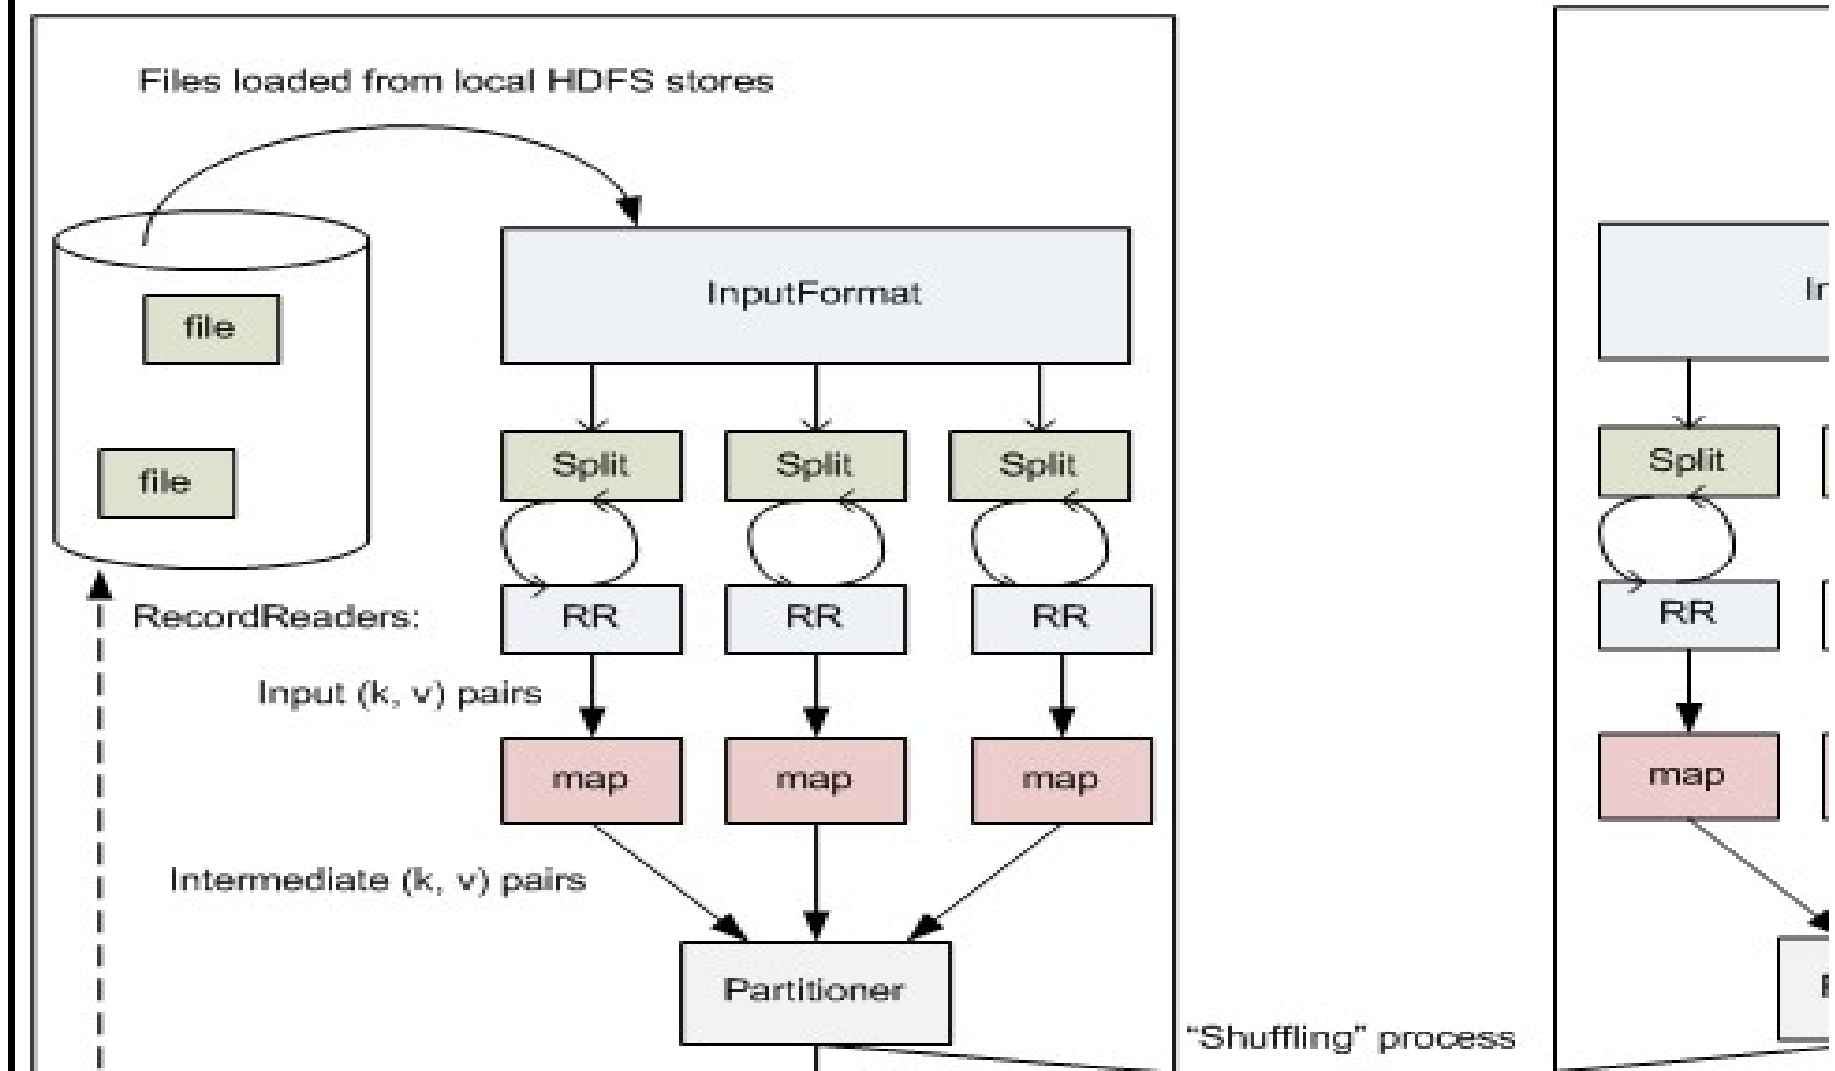
**Job** - A program is an execution of a Mapper and Reducer across a dataset.

**Task** - An execution of a Mapper or a Reducer on a slice of data.

**Task Attempt** - A particular instance of an attempt to execute a task on a SlaveNode.

# Data Flow

## Node 1



## Data Flow

**Input files:** This is where the data for a MapReduce task is initially stored.

- The format of these files is arbitrary.
- It is typical to be very large -- tens of gigabytes or more.
- **InputFormat:** It defines how these input files are split up and read.
- InputFormat is a class that provides the following functionality:
  - Selects the files or other objects that should be used for input
  - Defines the *InputSplits* that break a file into tasks
  - Provides a factory for *RecordReader* objects that read the file
- Several InputFormats are provided with Hadoop.
- An abstract type is called *FileInputFormat*.
- all InputFormats that operate on files inherit functionality and properties from this class.
- When starting a Hadoop job, FileInputFormat is provided with a path containing files to read.

## Data Flow

- The `FileInputFormat` will read all files in directory.
- Then it divides these files into one or more `InputSplits` each.
- We can choose which `InputFormat` to apply to our input files for a job by calling the `setInputFormat()` method of the *JobConf* object that defines the job.

The different `InputFormat` are:

- ***TextInputFormat***: It is default format.
  - It treats each line of each input file as a separate record.
  - This is useful for unformatted data or line-based records like log files.
- ***KeyValueInputFormat***: This format treats each line of input as a separate record.
- While the `TextInputFormat` treats the entire line as the value, the `KeyValueInputFormat` breaks the line itself into the key and value by searching for a tab character.

## Data Flow

- This is particularly useful for reading the output of one MapReduce job as the input to another.
- *SequenceFileInputFormat*: Reads special binary files that are specific to Hadoop.
- These files include many features designed to allow data to be rapidly read into Hadoop mappers.
- Sequence files are block-compressed and provide direct serialization and deserialization of several arbitrary data types (not just text).
- Sequence files can be generated as the output of other MapReduce tasks and are an efficient intermediate representation for data that is passing from one MapReduce job to another.

## Data Flow

- **RecordReader:** The InputSplit has defined a slice of work, but does not describe how to access it.
- The *RecordReader* class actually loads the data from its source and converts it into (key, value) pairs suitable for reading by the Mapper.
- The RecordReader instance is defined by the InputFormat.
- The default InputFormat, *TextInputFormat*, provides a *LineRecordReader*, which treats each line of the input file as a new value.
- The key associated with each line is its byte offset in the file.
- The RecordReader is invoked repeatedly on the input until the entire InputSplit has been consumed.
- Each invocation of the RecordReader leads to another call to the `map()` method of the Mapper.

## BASIC OPERATIONS

- Splits
- Mapper
- Portioning
- Sort
- shuffle
- Combiner
- Reducer

- **Splits:** Hadoop divides the input to a MapReduce job into fixed-size pieces called input splits, or just splits.
- Hadoop creates one map task for each split, which runs the userdefined map function for each record in the split.
- So if we are processing the splits in parallel, the processing is better load-balanced if the splits are small.




## BASIC OPERATIONS

- An InputSplit describes a unit of work that comprises a single map task in a MapReduce program.
- A MapReduce program applied to a data set, collectively referred to as a Job, is made up of several (possibly several hundred) tasks.
- Map tasks may involve reading a whole file, they involve reading only part of a file.

## BASIC OPERATIONS

- **Mapper:** The first phase of a MapReduce program is called *mapping*.
- A list of data elements are provided, one at a time, to a function called the *Mapper*, which transforms each element individually to an output data element.



## BASIC OPERATIONS

- The Mapper performs the user-defined work of the first phase of the MapReduce program.
- Given a key and a value, the `map()` method emits (key, value) pair(s) which are forwarded to the Reducers.
- A new instance of Mapper is instantiated in a separate Java process for each map task (InputSplit) that makes up part of the total job input.
- The individual mappers are intentionally not provided with a mechanism to communicate with one another in any way.
- OutputCollector class
  - `collect()`
- Reporter class
  - `getInputSplit()`
  - `setStatus()`
  - `incrCounter()`

## BASIC OPERATIONS

- **Partitioner**: A partitioner works like a condition in processing an input dataset.
- The partition phase takes place after the Map phase and before the Reduce phase.
- The number of partitioners is equal to the number of reducers.
- That means a partitioner will divide the data according to the number of reducers.
- Therefore, the data passed from a single partitioner is processed by a single Reducer.
- A partitioner partitions the key-value pairs of intermediate Map-outputs.
- It partitions the data using a user-defined condition, which works like a hash function.
- The total number of partitions is same as the number of Reducer tasks for the job.

## BASIC OPERATIONS

- *Partitioner* interface

- Syntax:

```
public interface Partitioner<K, V> extends JobConfigurable  
{  
    int getPartition(K key, V value, int numPartitions);  
}
```

- *HashPartitioner* (default partitioner)

- `hashCode()` modulo the number of partitions total to determine which partition to send a given (key, value) pair to.

## BASIC OPERATIONS

- **Shuffle:**

- **Sort:** **Context** class (user-defined class) collects the matching valued keys as a collection.

**RawComparator** class to sort the key-value pairs for similar key-value pair.

Example: (K2, {V2, V2, ...})

**Combiner:** The Combiner is a "mini-reduce" process which operates only on data generated by one machine.

- The Combiner class is used in between the Map class and the Reduce class to reduce the volume of data transfer between Map and Reduce.

- A combiner does not have a predefined interface and it must implement the Reducer interface's reduce() method.

*conf.setCombinerClass(Reduce.class);*

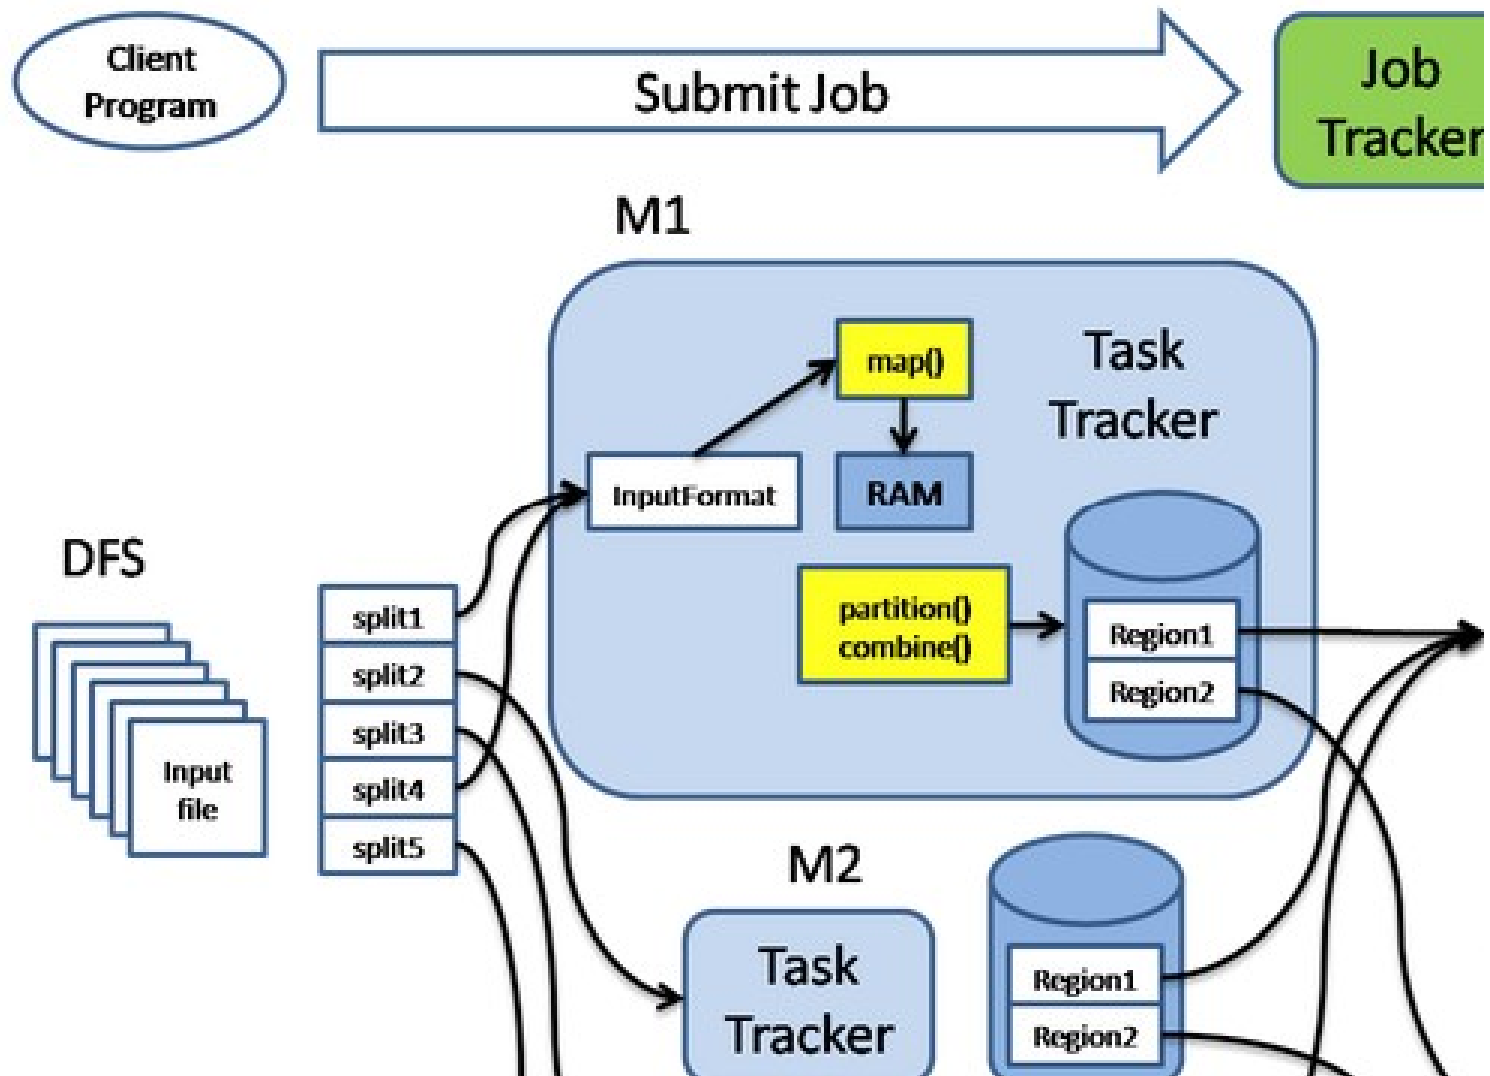
## BASIC OPERATIONS

• **Reduce:** For each key in the partition assigned to a Reducer.

- *Reducer* class
- *JobContext.getConfiguration()* is used to configure
- *reduce()*

**reduce**(KEYIN key, Iterable<VALUEIN> values,  
org.apache.hadoop.mapreduce.Reducer.Context  
context)

# How MapReduce works

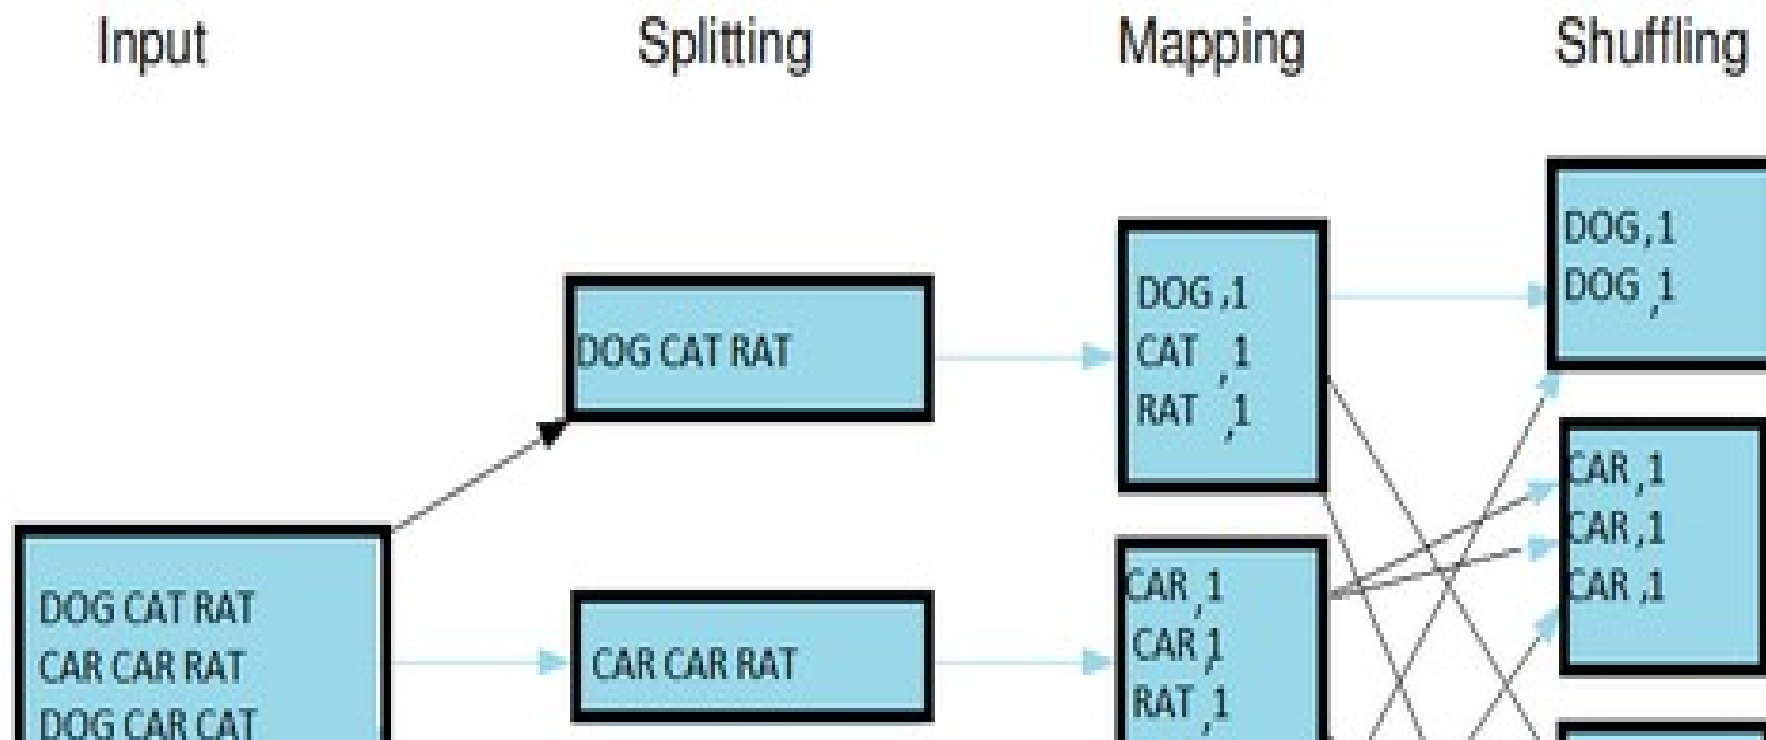




# Word count Example

Count the number of words in the file.

The overall MapReduce word count process



# Mapreduce life Cycle

## Job Submission

The `submit()` method on job creates an internal instance of `JobSubmitter` and calls `submitJobInternal()` method on it.

- `waitForCompletion()` polls the job's progress once a second after submitting the job.

On calling this method following happens.

- It goes to JobTracker and gets a jobId for the job
- Perform checks if the the output directory has been specified or not.
- If specified , whether the directory already exists or is new.
- And throws error if any such thing fails.
- Computes input split and throws error if it fails to do so, because the input paths don't exist.
- Copies the resources to JobTracker file system in a directory named after Job Id.
- These resources include configuration files, job jar file, and computed input splits.

Finally it calls `submitJob()` method on JobTracker.

# Mapreduce life Cycle

## Job Initialization

Job tracker performs following steps

- Creates book keeping object to track tasks and their progress

For each input split creates a map tasks.

- The number of reduce tasks is defined by the configuration `mapred.reduce.tasks` set by `setNumReduceTasks()`.

- Tasks are assigned taskId's at this point.

- In addition to this 2 other tasks are created: `Job initialization task` and `Job clean up task`, and are run by tasktrackers

- `FileOutputCommitter` creates the output directory to store the tasks output as well as temporary output

- Job clean up tasks which delete the temporary directory after the job is complete.

# Mapreduce life Cycle

## Task Assignment

TaskTracker sends a heartbeat to jobtracker every five seconds. It will indicate whether it is ready to run a new task. They also send the available slots on them.

Here is how job allocation takes place.

- JobTracker first selects a job to select the task from, based on job scheduling algorithms.
- The default scheduler fills empty map task before reduce task slots.
- For a map task then it chooses a tasktracker which is in following order of priority: **data-local, rack-local and then network**.
- For a reduce task, it simply chooses a task tracker which has empty slots.

# Mapreduce life Cycle

## Task Execution:

- TaskTracker copies the job jar file from the shared filesystem (HDFS).
- Tasktracker creates a local working directory, and un-jars the jar file into the local file system.
- It then creates an instance of TaskRunner

## Mapreduce life Cycle

### Action:

Tasktracker starts TaskRunner in a new JVM to run the map or reduce task.

- Seperate process is needed so that the TaskTracker does not crash in case of bug in user code or JVM.
- The child process communicates its progress to parent process .
- Each task can perform setup and cleanup actions, which are run in the same JVM as the task itself, based on **OutputComitter**.

# Mapreduce life Cycle

## Job/Task Progress

Here is how the progress is monitored of a job/task

- JobClient keeps polling the JobTracker for progress.
- Each child process reports its progress to parent task tracker.
- If a task reports progress, it sets a flag to indicate that the status change should be sent to the tasktracker.
- The flag is checked in a separate thread every 3 seconds, and if set it notifies the tasktracker of the current task status.
- Task tracker sends its progress to JobTracker over the heartbeat for every five seconds.
- JobTracker then assembles task progress from all task trackers and keeps a view of job.
- The Job receives the latest status by polling the jobtracker every second.

# Mapreduce life Cycle

## Job Completion

On Job Completion the clean up task is run.

- Task sends the task tracker job completion. Which is sent to job tracker.
- Job Tracker then send the job completion message to client.
- Jobtracker cleans up its working state for the job and instructs tasktrackers to do the same, It cleans up all the temporary directories.
- This causes jobclient's `waitForJobToComplete()` method to return.



## Mapreduce Program

```
package hadoop;  
import java.util.*;  
import java.io.IOException;  
import java.io.IOException;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.conf.*;  
import org.apache.hadoop.io.*;  
import org.apache.hadoop.mapred.*;  
import org.apache.hadoop.util.*;
```

# Mapreduce Program

```
public class ProcessUnits
{
    public static class E_EMapper extends MapReduceBase implements
    Mapper<LongWritable,                               /*Input key Type */
    Text,                                               /*Input value Type*/
    Text,                                               /*Output key Type*/
    IntWritable>                                       /*Output value Type*/
    {
        //Map function
        public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output,
        Reporter reporter) throws IOException
        {
            String line = value.toString();
            String lasttoken = null;
            StringTokenizer s = new StringTokenizer(line, "\t");
            String year = s.nextToken();
            while(s.hasMoreTokens())
            {
                lasttoken=s.nextToken();
            }
        }
    }
}
```

# Mapreduce Program

```
int avgprice = Integer.parseInt(lasttoken);  
output.collect(new Text(year), new IntWritable(avgprice));  
}  
}
```

# Mapreduce Program

## Reducer class

```
public static class E_EReduce extends MapReduceBase implements
    Reducer< Text, IntWritable, Text, IntWritable >
{
    //Reduce function
    public void reduce(Text key, Iterator <IntWritable> values, OutputCollector<Text,
    IntWritable> output, Reporter reporter) throws IOException
    {
        int maxavg=30;
        int val=Integer.MIN_VALUE;
        while (values.hasNext())
        {
            if((val=values.next().get())>maxavg)
            {
                output.collect(key, new IntWritable(val));
            }
        }
    }
}
```

# Mapreduce Program

```
//Main function
public static void main(String args[])throws Exception
{
    JobConf conf = new JobConf(Eleunits.class);
    conf.setJobName("max_eletricityunits");
    conf.setOutputKeyClass(Text.class);
    conf.setOutputValueClass(IntWritable.class);
    conf.setMapperClass(E_EMapper.class);
    conf.setCombinerClass(E_EReduce.class);
    conf.setReducerClass(E_EReduce.class);
    conf.setInputFormat(TextInputFormat.class);
    conf.setOutputFormat(TextOutputFormat.class);
    FileInputFormat.setInputPaths(conf, new Path(args[0]));
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));
    JobClient.runJob(conf);
}
}
```

## Mapreduce Program

<b>1979</b>	23	23	2	43	24	25	26	26	:
<b>1980</b>	26	27	28	28	28	30	31	31	:
<b>1981</b>	31	32	32	32	33	34	35	36	:

The above data is saved as **sample.txt**

# Mapreduce Program

Step-1: create a directory to store the compiled java classes.

```
$ mkdir unit
```

Step-2: Download Hadoop-core-1.2.1.jar, which is used to compile and execute the MapReduce program.

Step 3: compile the **ProcessUnits.java** program and to create a jar for the program.

```
$ javac -classpath hadoop-core-1.2.1.jar -d units ProcessUnits.java
```

```
$ jar -cvf units.jar -C units/
```

Step 4: create an input directory in HDFS.

```
$HADOOP_HOME/bin/hadoop fs -mkdir input_dir
```

Step 5: copy the input file named **sample.txt** in the input directory of HDFS.

```
$HADOOP_HOME/bin/hadoop fs -put /home/hadoop/sample.txt input_dir
```

Step 6: verify the files in the input directory

```
$HADOOP_HOME/bin/hadoop fs -ls input_dir/
```

Step 7: run the application

```
$HADOOP_HOME/bin/hadoop jar units.jar hadoop.ProcessUnits input_dir  
output_dir
```

# Mapreduce Program

```
INFO mapreduce.Job: Job job_1414748220717_0002
completed successfully
14/10/31 06:02:52
INFO mapreduce.Job: Counters: 49
```

## File System Counters

```
FILE: Number of bytes read=61
FILE: Number of bytes written=279400
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0

HDFS: Number of bytes read=546
HDFS: Number of bytes written=40
HDFS: Number of read operations=9
HDFS: Number of large read operations=0
HDFS: Number of write operations=2 Job Counters
```



# Mapreduce Program

Step-8: verify the resultant files in the output folder.

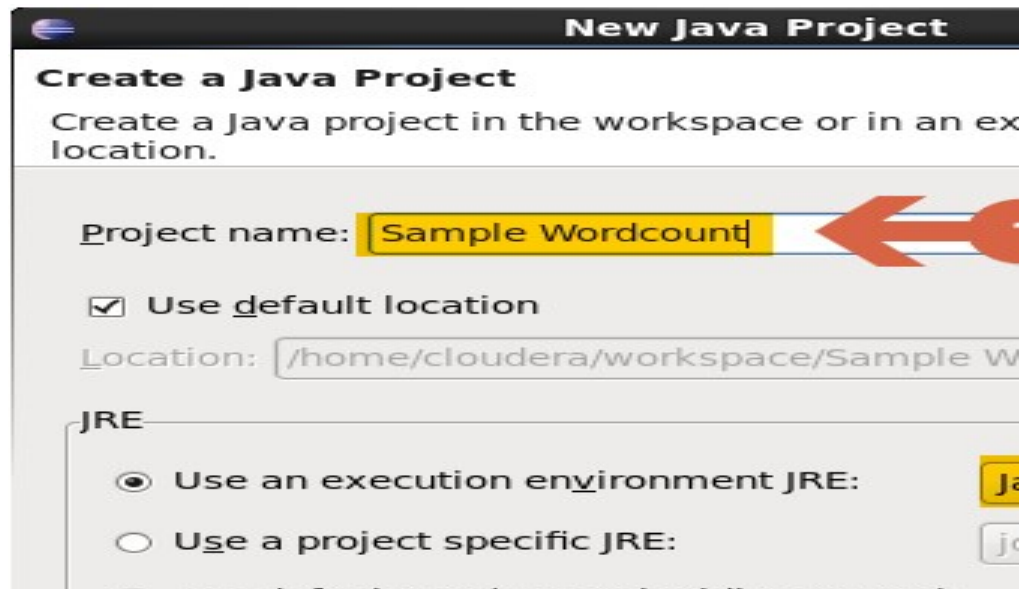
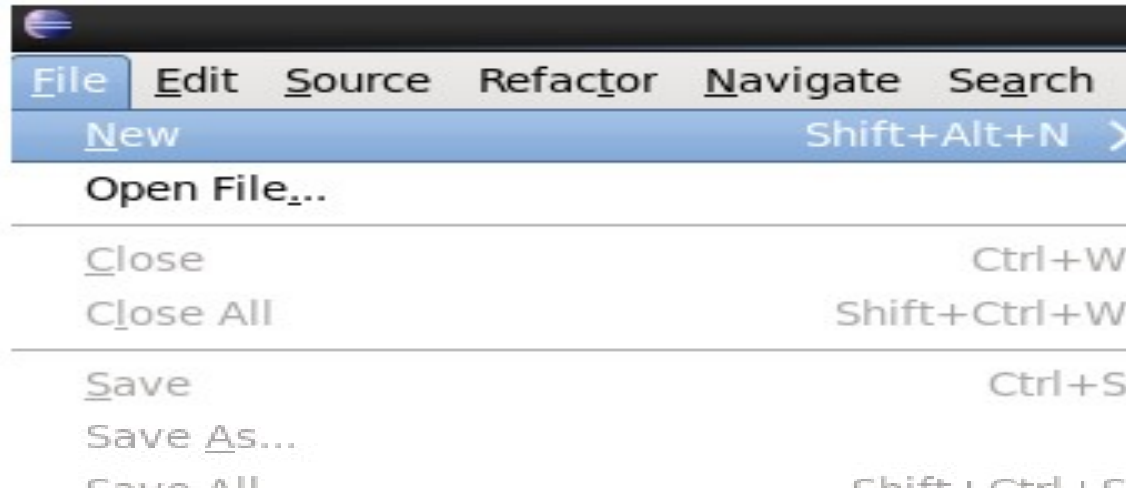
```
$HADOOP_HOME/bin/hadoop fs -ls output_dir/
```

Step-9: see the output

```
$HADOOP_HOME/bin/hadoop fs -cat output_dir
```

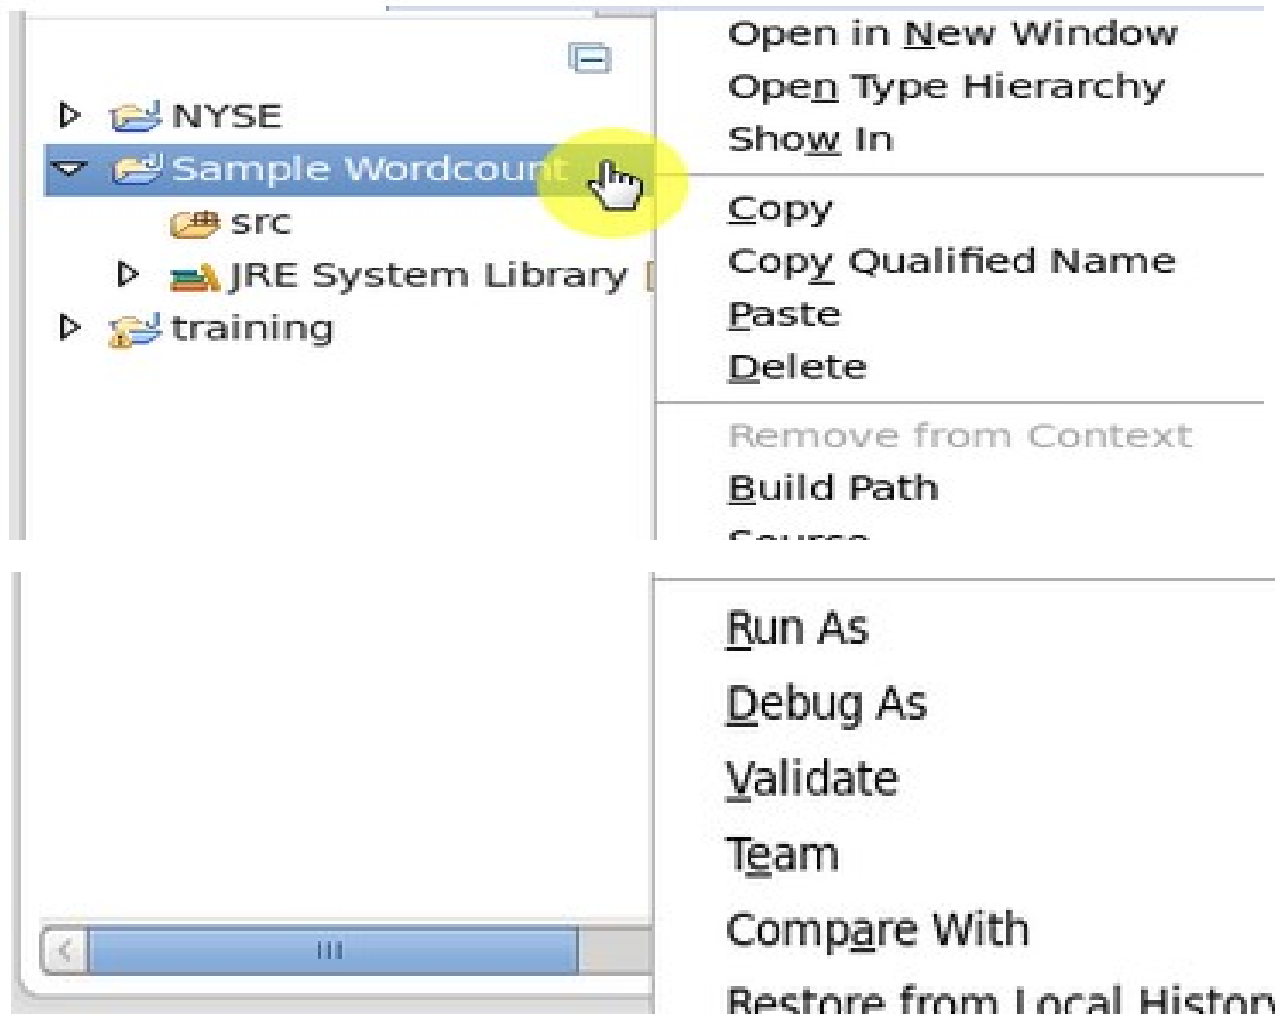
# Wordcount Program

Step-1: Open Eclipse> File > New > Java Project >( Name it – Sample Wordcount)  
> Finish

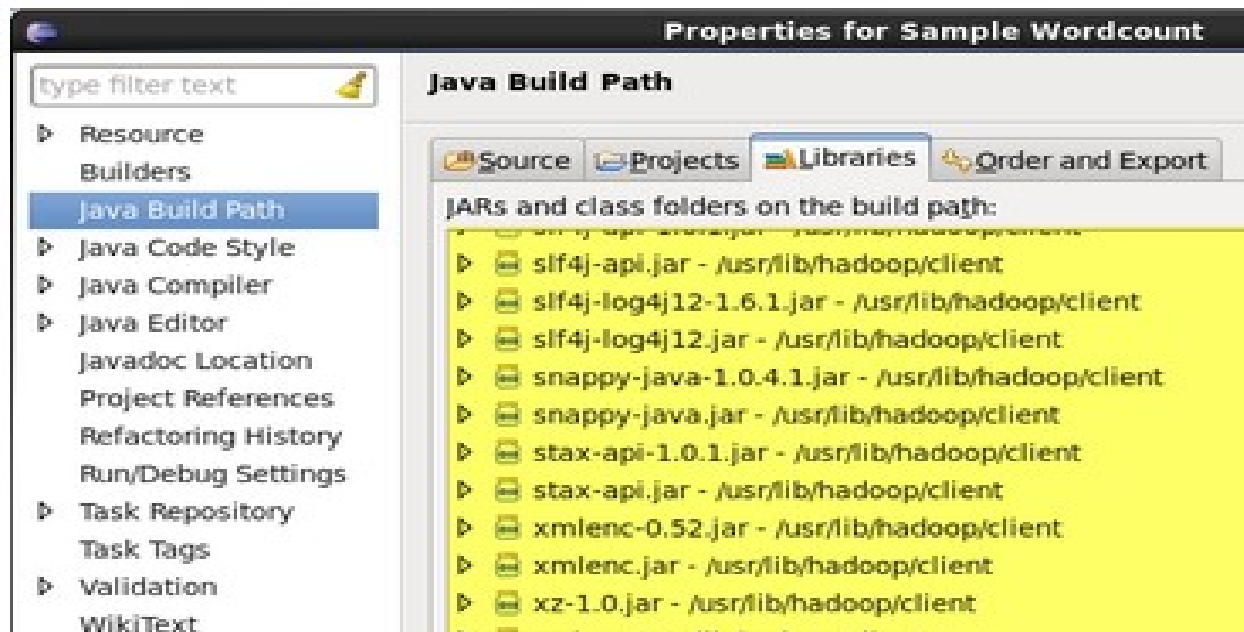
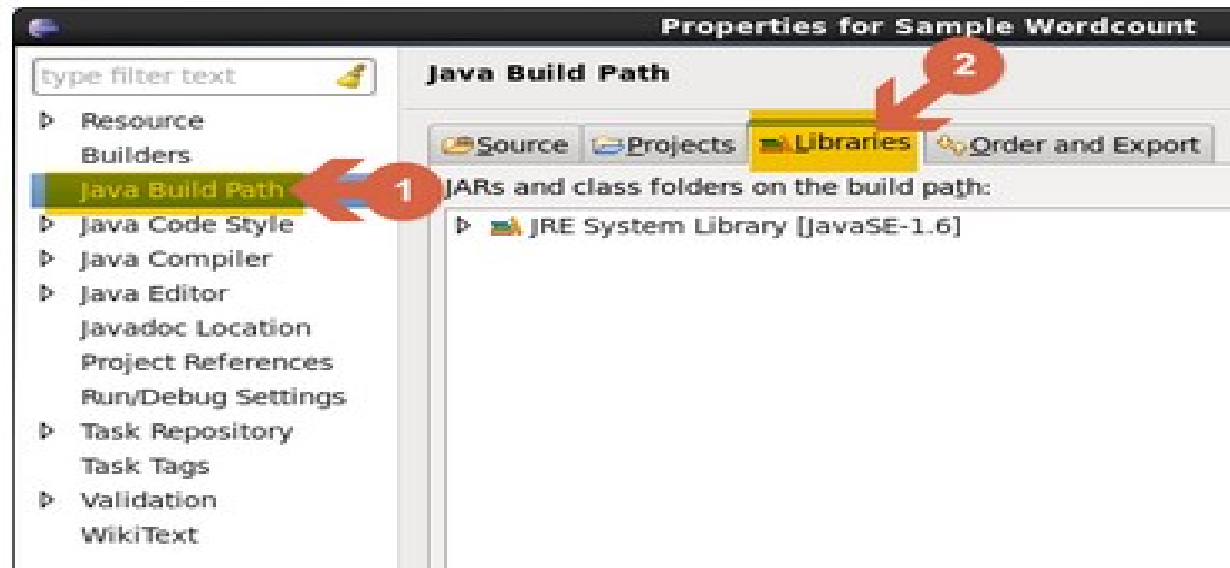


# Wordcount Program

Step-2: Get references to hadoop libraries by clicking on Add JARS

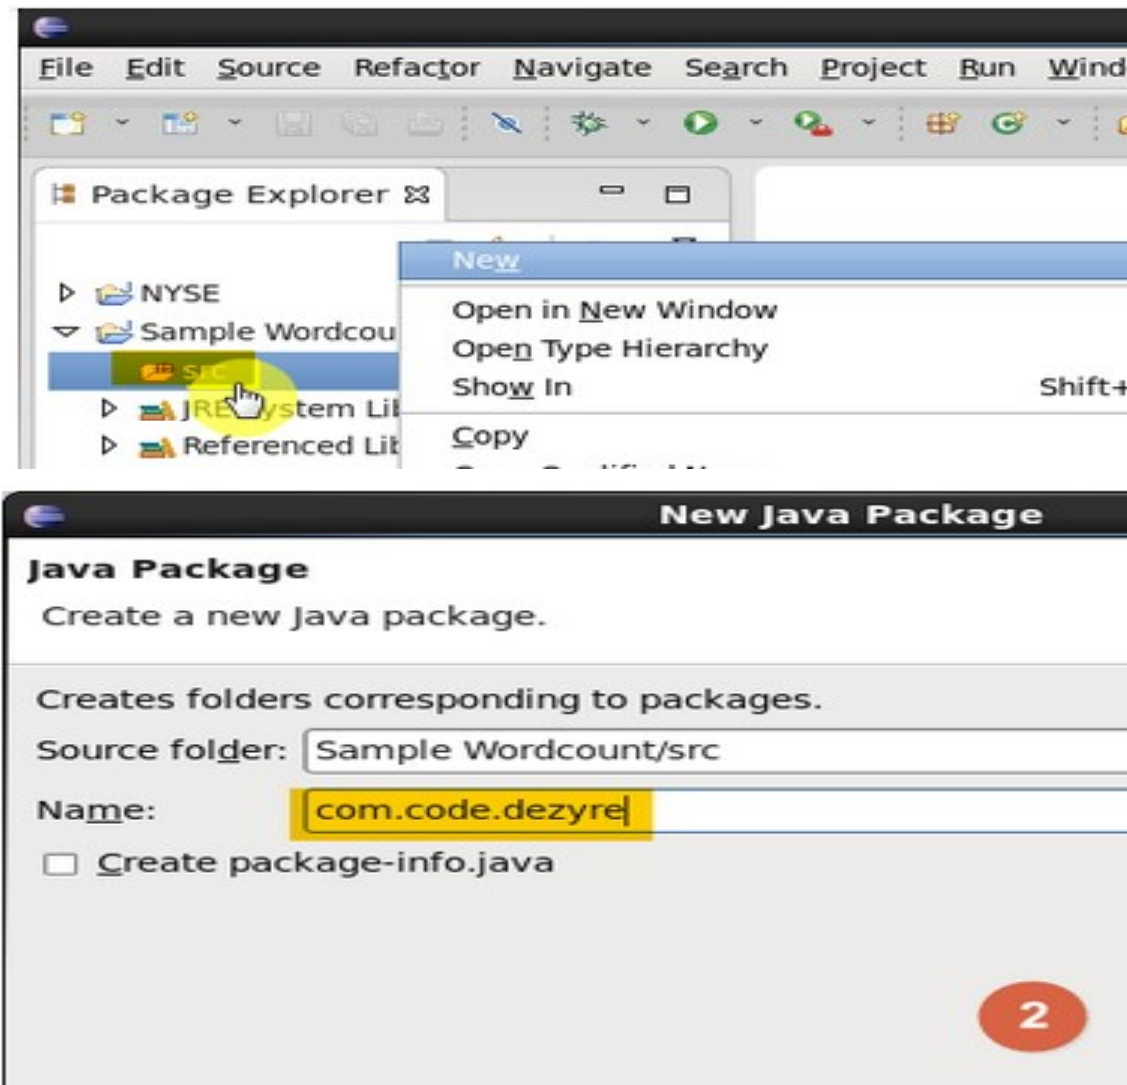


# Wordcount Program



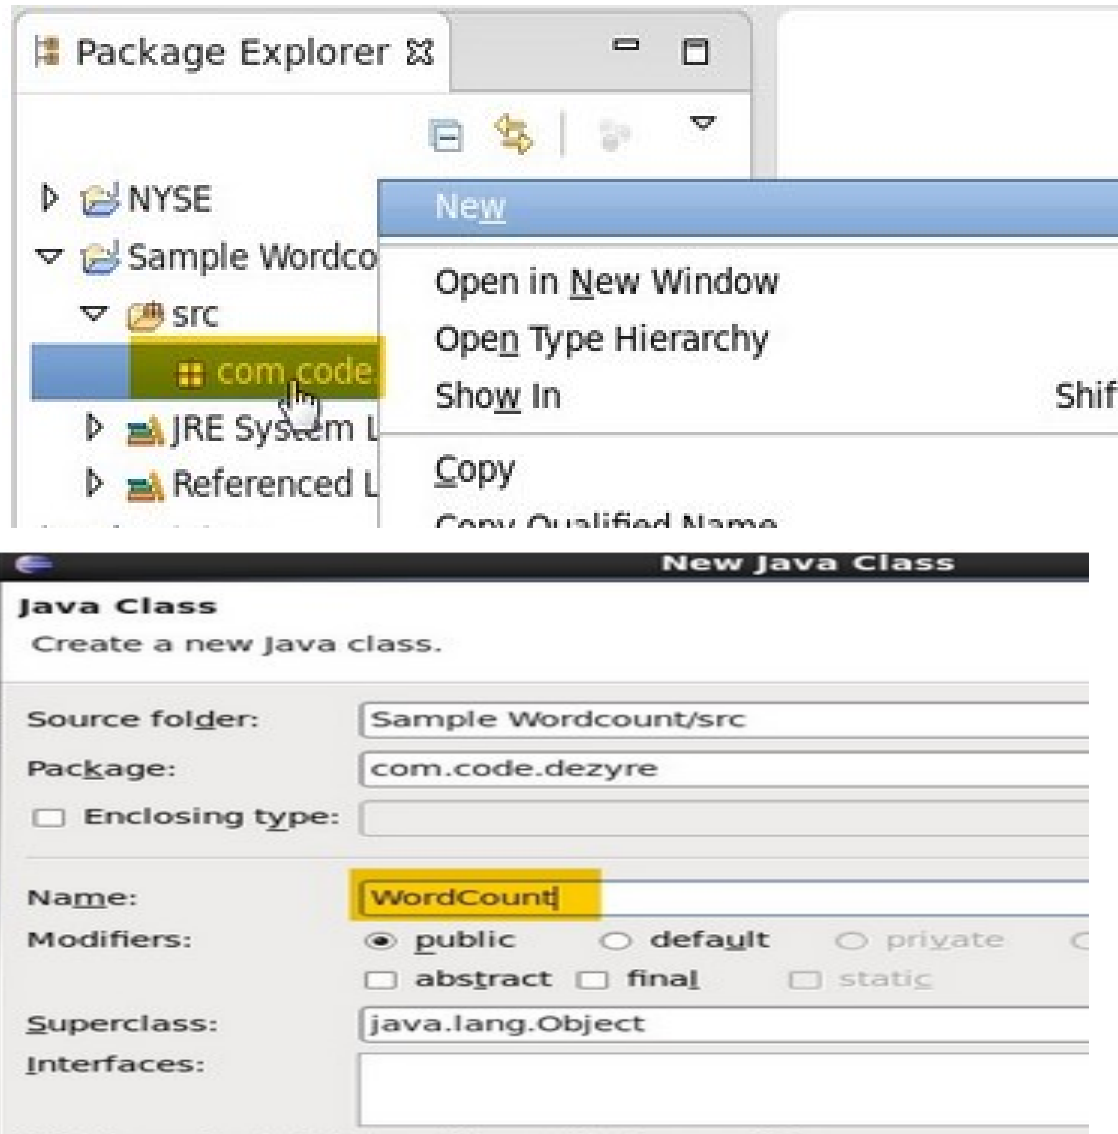
# Wordcount Program

Step-3: Create a new package within the project Right Click > New > Package (Name it) > Finish



# Wordcount Program

Step-4: Create a WordCount class under the project

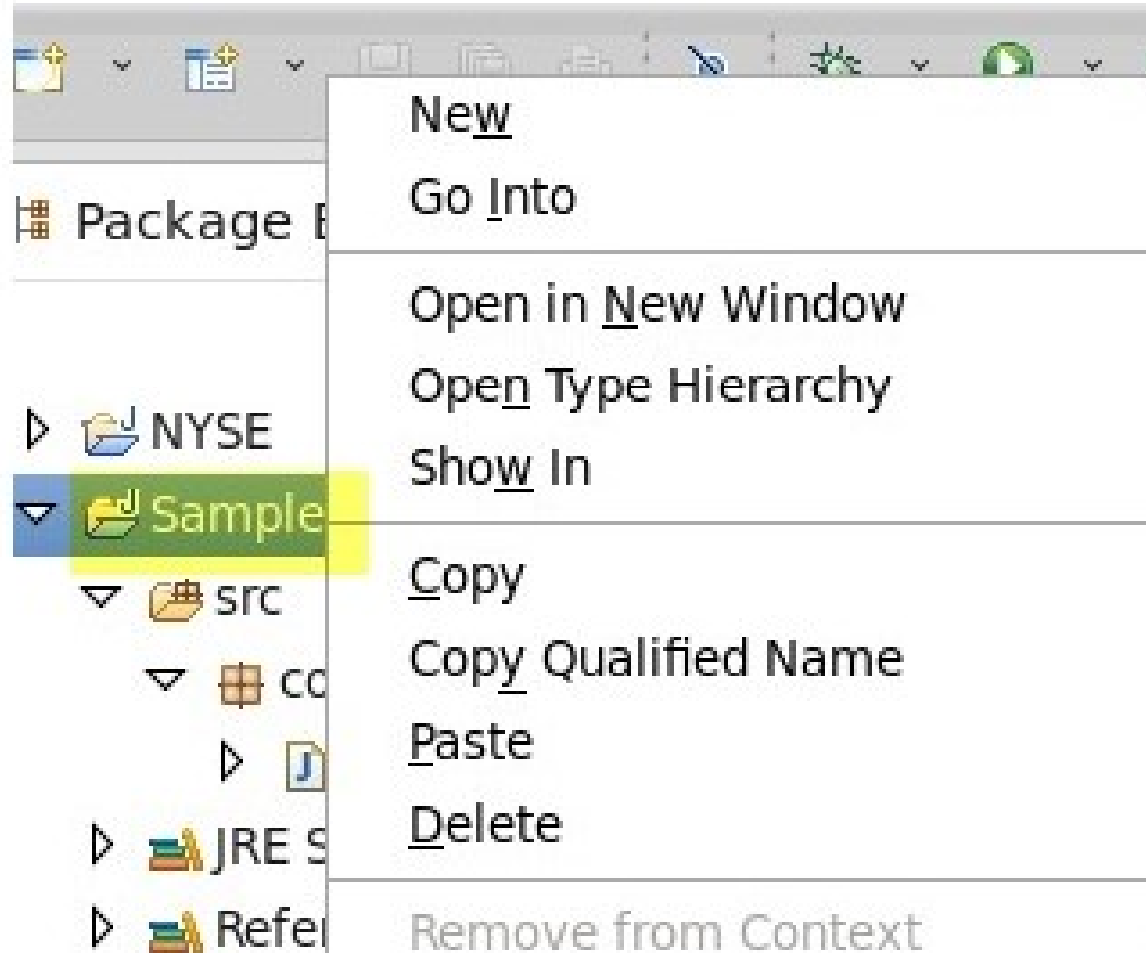


# Wordcount Program

Step-5: Write mapper code, reducer code and configuration code

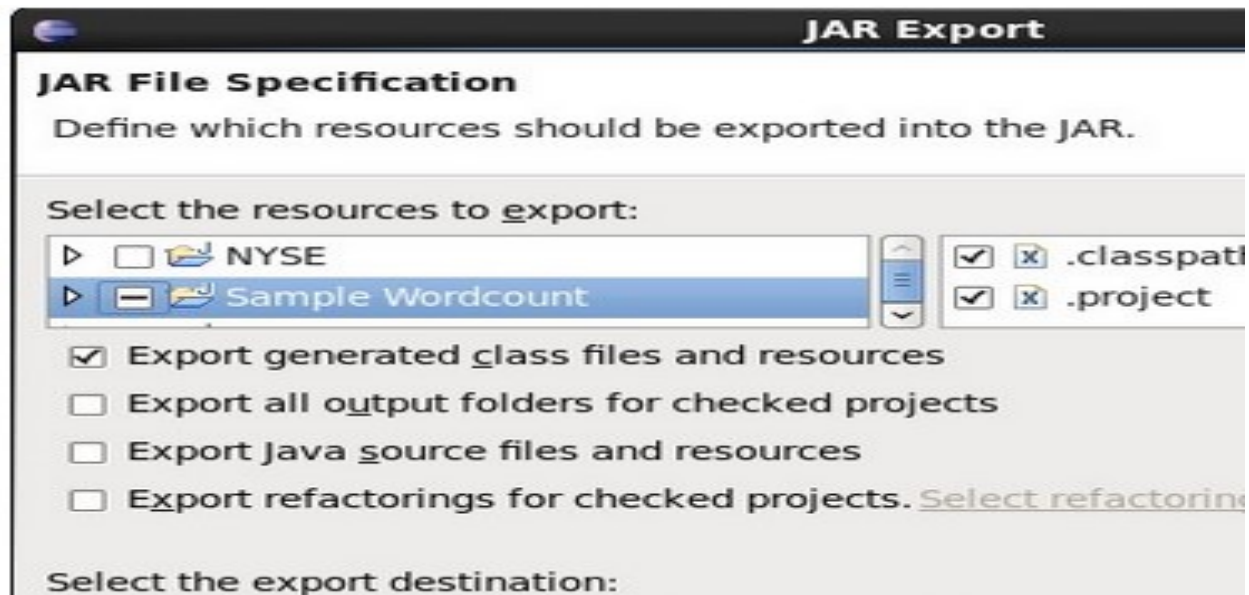
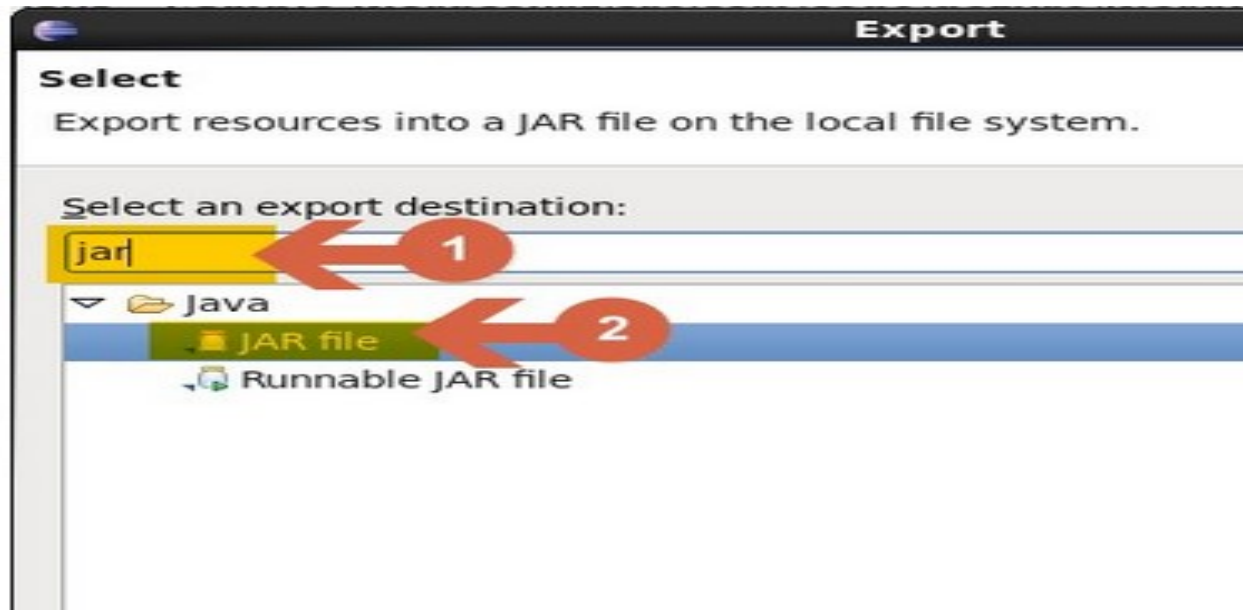
# Wordcount Program

Step-6: Create the JAR file for the wordcount class





# Wordcount Program




# Wordcount Program



# Hadoop InputOutput Format

## Input Format



# Hadoop InputOutput Format

The InputFormat class is one of the fundamental classes in the Hadoop MapReduce framework which provides the following functionality:

- The files or other objects that should be used for input is selected by the InputFormat.
- InputFormat defines the Data splits, which defines both the size of individual Map tasks and its potential execution server.
- InputFormat defines the RecordReader, which is responsible for reading actual records from the input files.

- FileInputFormat in Hadoop**
- **TextInputFormat**
- KeyValueTextInputFormat**
- SequenceFileInputFormat**
- SequenceFileAsTextInputFormat**
- SequenceFileAsBinaryInputFormat**
- NLineInputFormat**
- **DBInputFormat.**

## How we get the data to mapper?

2 methods used to get the data to mapper in [MapReduce](#):

- getsplits()
- createRecordReader()

```
public abstract class InputFormat<K, V>
{
    public abstract List<InputSplit> getSplits(JobContext context)
    throws IOException, InterruptedException;
    public abstract RecordReader<K, V>
    createRecordReader(InputSplit split,
    TaskAttemptContext context) throws IOException,
    InterruptedException;
}
```

# Hadoop Input Format

**FileInputFormat in Hadoop:** FileInputFormat in Hadoop is the base class for all file-based InputFormats.

- FileInputFormat specifies input directory where data files are located.
- When we start a Hadoop job, FileInputFormat is provided with a path containing files to read.
- FileInputFormat will read all files and divides these files into one or more InputSplits.
- TextInputFormat:** TextInputFormat in Hadoop is the default InputFormat of MapReduce.
- TextInputFormat treats each line of each input file as a separate record and performs no parsing.

**KeyValueTextInputFormat:** KeyValueTextInputFormat in Hadoop is similar to TextInputFormat as it also treats each line of input as a separate record.

- While TextInputFormat treats entire line as the value, but the KeyValueTextInputFormat breaks the line itself into key and value by a tab character ('\t').

**SequenceFileInputFormat:** SequenceFileInputFormat in Hadoop is an InputFormat which reads sequence files.

- Sequence files are binary files that stores sequences of binary key-value pairs.
- Sequence files are block-compressed and provide direct serialization and deserialization of several arbitrary data types (not just text).
- Key & Value both are user-defined.

# Hadoop Input Format

**SequenceFileAsTextInputFormat:** SequenceFileAsTextInputFormat in Hadoop is another form of SequenceFileInputFormat which converts the sequence file key values to Text objects.

- By calling 'toString()' conversion is performed on the keys and values.
- This InputFormat makes sequence files suitable input for streaming.

**SequenceFileAsBinaryInputFormat:** SequenceFileAsBinaryInputFormat in Hadoop is a SequenceFileInputFormat using which we can extract the sequence file's keys and values as a binary object.

**NLineInputFormat:** It is another form of TextInputFormat where the keys are byte offset of the line and values are contents of the line.

- Each mapper receives a variable number of lines of input with TextInputFormat and KeyValueTextInputFormat.
- The number depends on the size of the split and the length of the lines.
- And if we want our mapper to receive a fixed number of lines of input, then we use NLineInputFormat.

**Example:** N is the number of lines of input that each mapper receives. By default (N=1), each mapper receives exactly one line of input.

- If N=2, then each split contains two lines. One mapper will receive the first two Key-Value pairs and another mapper will receive the second two key-value pairs.

# Hadoop Input Format

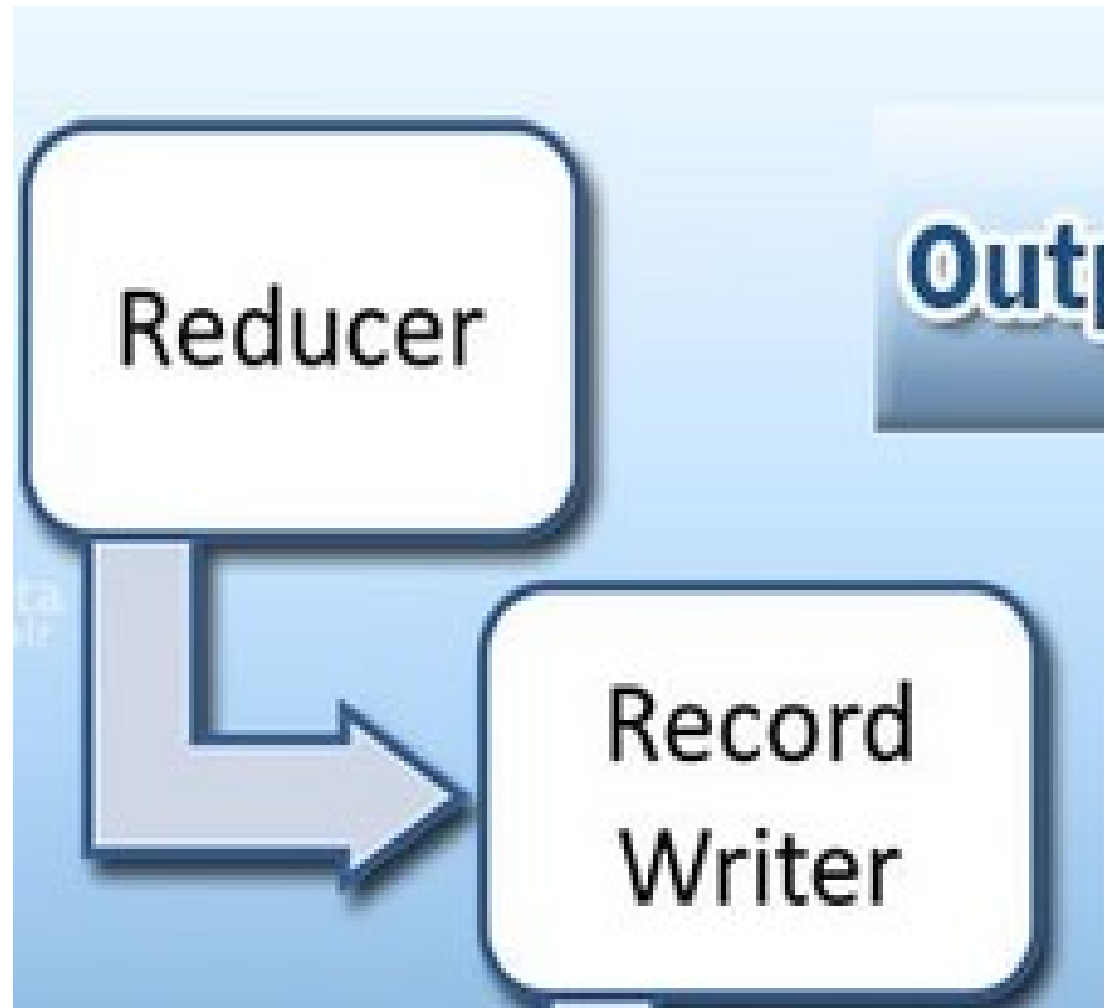
**DBInputFormat:** DBInputFormat in Hadoop is an InputFormat that reads data from a relational database, using JDBC.

- it doesn't have portioning capabilities.
- so we need to be careful to read to many mappers.
- joining is used to make datasets from [HDFS](#) using MultipleInputs.
- Here Key is LongWritable while Value is DBWritable.



# Hadoop InputOutput Format

## Output Format



# Hadoop Output Format

OutputFormat is used to write to files on the local disk or in HDFS.

- OutputFormat describes the output-specification for a Map-Reduce job.
- Based on output specification following things happened:
  - MapReduce job checks that the output directory.
  - OutputFormat provides the RecordWriter implementation to be used to write the output files of the job.
  - Output files are stored in a FileSystem.
  - FileOutputFormat.setOutputPath() method is used to set the output directory.
  - Every Reducer writes a separate file in a common output directory.

# Types of Hadoop Output Format

- `textOutputFormat`,
- `sequenceFileOutputFormat`,
- `mapFileOutputFormat`,
- `sequenceFileAsBinaryOutputFormat`,
- `DBOutputFormat`,
- `LazyOutputForma`, and
- `MultipleOutputs`

## **TextOutputFormat:**

- MapReduce default OutputFormat is TextOutputFormat, it writes (key, value) pairs on individual lines of text files
- The keys and values can be of any type.
- TextOutputFormat turns it to string by calling `toString()`.
- Each key-value pair is separated by a tab character.
- It can be changed using `MapReduce.output.textoutputformat.separator` property.
- `KeyValueTextOutputFormat` is used for reading output text files since it breaks lines into key-value pairs based on a configurable separator.

# MapReduce Join

- It is used to achieve larger dataset.
- Joining of two datasets begin by comparing size of each dataset.
- If one dataset is smaller as compared to the other dataset then smaller dataset is distributed to every datanode in the cluster.
- Once it is distributed, either Mapper or Reducer uses smaller dataset to perform lookup for matching records from large dataset and then combine those records to form output records.

**Map-side join** - When the join is performed by the mapper, it is called as map-side join.

- In this type, the join is performed before data is actually consumed by the map function.
- It is mandatory that the input to each map is in the form of a partition and is in sorted order.
- Also, there must be an equal number of partitions and it must be sorted by the join key.

**Reduce-side join** - When the join is performed by the reducer, it is called as reduce-side join.

- There is no necessity in this join to have dataset in a structured form (or partitioned).
- Here, map side processing emits join key and corresponding tuples of both the tables.
- As an effect of this processing, all the tuples with same join key fall into the same reducer which then joins the records with same join key.

# MapReduce Join Example

<u>Name</u>	<u>Salary</u>	<u>Dept_ID</u>
Sumit	700000	5
Dilip	750000	2

Dept_ID	Department
2	Marketing Department
5	Finance


# MapReduce Join Example-1

<u>Name</u>	<u>Salary</u>	<u>Dept_ID</u>
Sumit	700000	5
Dilip	750000	2
Amar	500000	5
Abhijit	800000	5


*Employee*

<u>Dept_ID</u>	<u>Name</u>
2	Marketing
5	Finance
3	Sales

*Department*

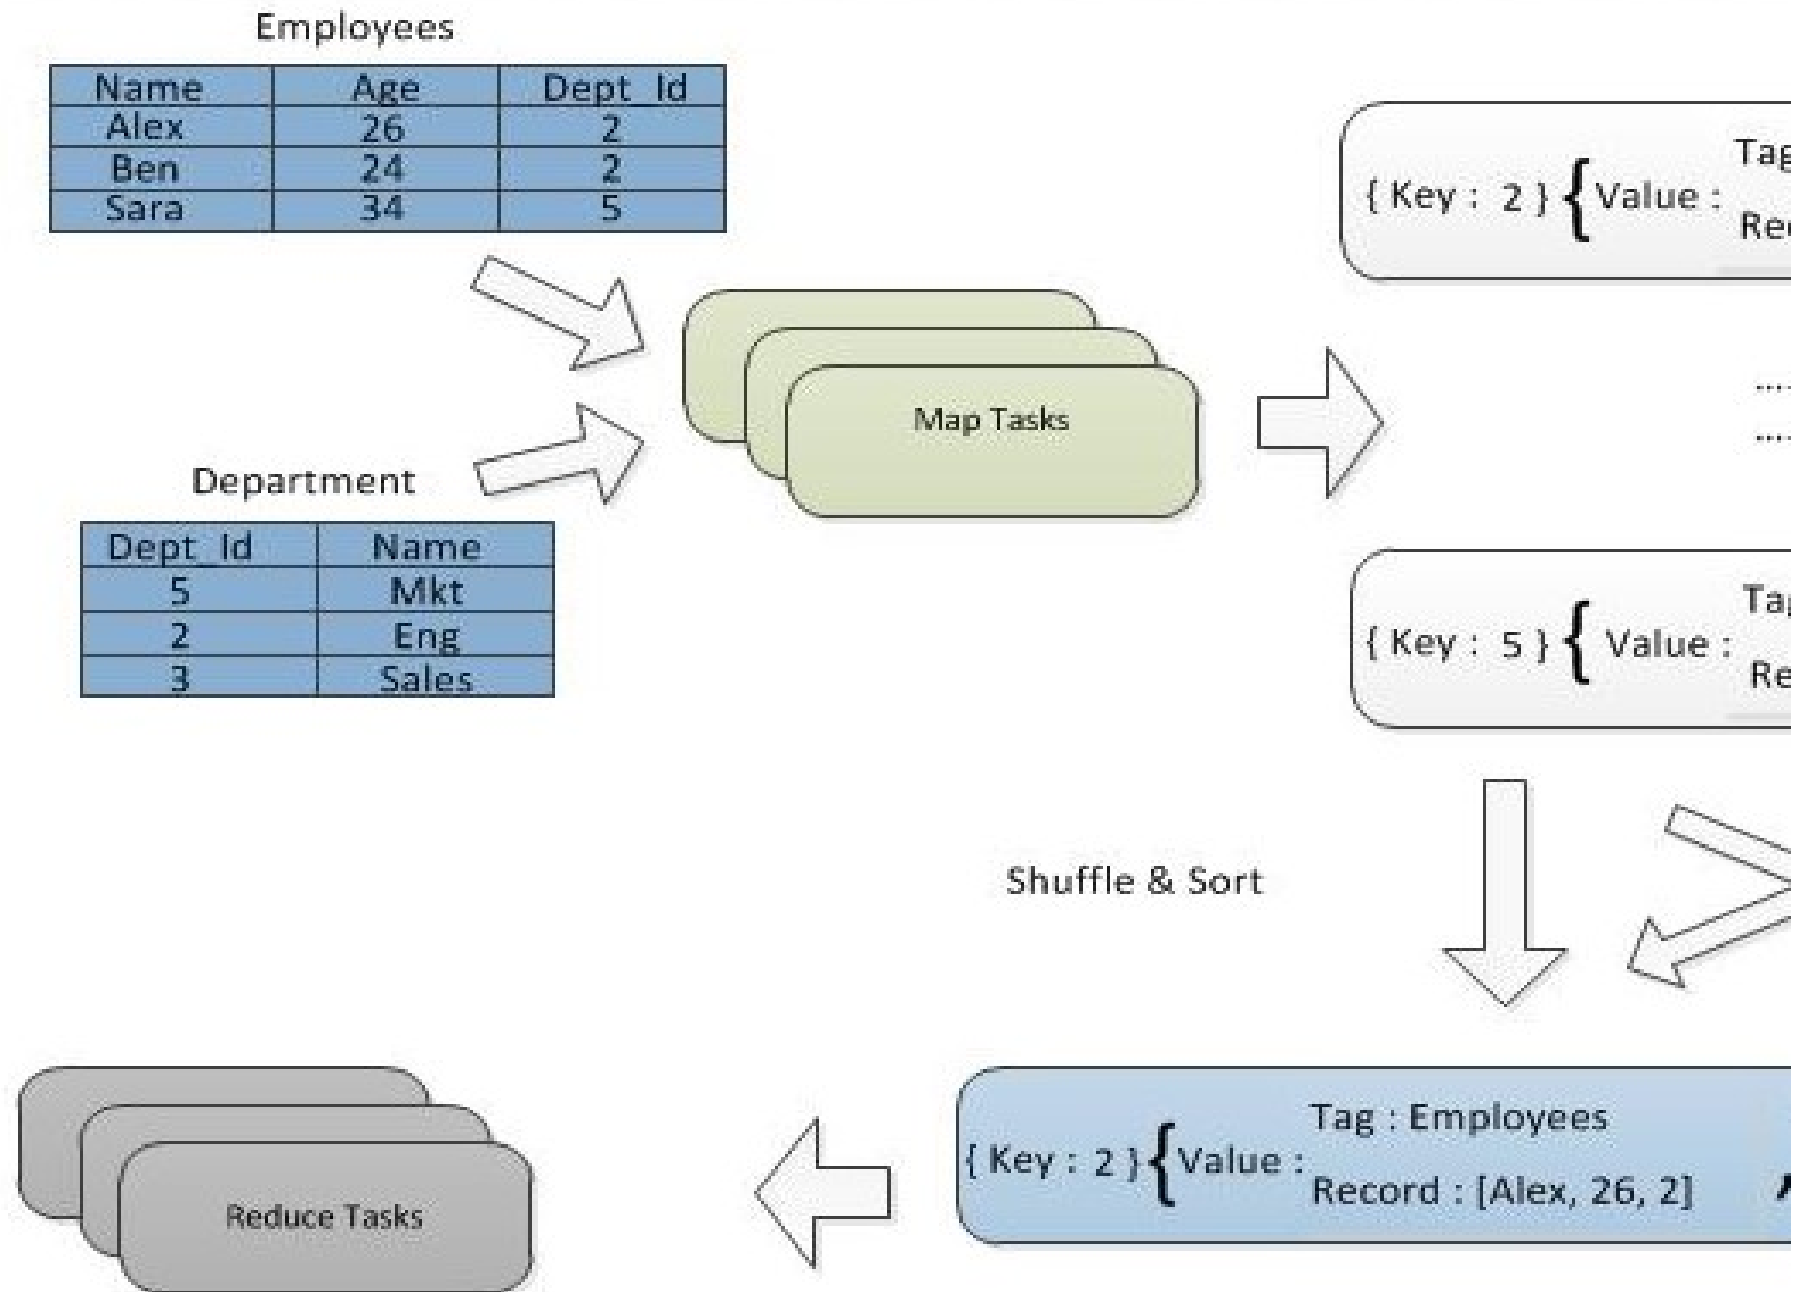


*Shuffle, Sort*



Key=2, {Value=(  
    {Value=(

## MapReduce Join Example-2



## pseudo code

```
map (K table, V rec)
{
    dept_id = rec.Dept_Id
    tagged_rec.tag = table
    tagged_rec.rec = rec
    emit(dept_id, tagged_rec)
}
```

```
reduce (K dept_id, list<tagged_rec> tagged_recs)
{
    for (tagged_rec : tagged_recs)
    {
        for (tagged_rec1 : tagged_recs)
        {
            if (tagged_rec.tag != tagged_rec1.tag)
            {
                joined_rec = join(tagged_rec, tagged_rec1)
            }
            emit (tagged_rec.rec.Dept_Id, joined_rec)
        }
    }
}
```



# NOW TIME FOR U !!!!!

- 1: How Hadoop is different from Data-mining and Data-ware Housing?
- 2: Explain with diagram, How Mapreduce works?
- 3: Explain Streaming and pipe.
- 4: What is counter? Explain different types of counter.
- 5: Briefly explain different types of sorting used in mapreduce.

THANKS